

Lesson 3 Exercise 1 Three Queries Three Tables

January 23, 2023

1 Lesson 3 Exercise 1: Three Queries Three Tables

1.0.1 Walk through the basics of creating a table in Apache Cassandra, inserting rows of data, and doing a simple CQL query to validate the information. You will practice Denormalization, and the concept of 1 table per query, which is an encouraged practice with Apache Cassandra.

1.0.2 Remember, replace ##### with your answer.

We will use a python wrapper/ python driver called `cassandra` to run the Apache Cassandra queries. This library should be preinstalled but in the future to install this library you can run this command in a notebook to install locally: `! pip install cassandra-driver #####` More documentation can be found here: <https://datastax.github.io/python-driver/>

Import Apache Cassandra python package

```
In [1]: import cassandra
```

1.0.3 Create a connection to the database

```
In [2]: from cassandra.cluster import Cluster
        try:
            cluster = Cluster(['127.0.0.1']) #If you have a locally installed Apache Cassandra i
            session = cluster.connect()
        except Exception as e:
            print(e)
```

1.0.4 Create a keyspace to work in

```
In [3]: try:
        session.execute("""
            CREATE KEYSPACE IF NOT EXISTS udacity
            WITH REPLICATION =
            { 'class' : 'SimpleStrategy', 'replication_factor' : 1 }""")
    except Exception as e:
        print(e)
```

Connect to our Keyspace. Compare this to how we had to create a new session in PostgreSQL.

```
In [4]: try:
        session.set_keyspace('udacity')
    except Exception as e:
        print(e)
```

1.0.5 Let's imagine we would like to start creating a Music Library of albums.

1.0.6 We want to ask 3 questions of the data

1. Give every album in the music library that was released in a given year `select * from music_library WHERE YEAR=1970 #####` **2. Give every album in the music library that was created by a given artist**
`select * from artist_library WHERE artist_name="The Beatles" #####` **3. Give all the information from the music library about a given album** `select * from album_library WHERE album_name="Close To You"`

1.0.7 Because we want to do three different queries, we will need different tables that partition the data differently.

1.0.8 TO-DO: Create the tables.

```
In [6]: query = "CREATE TABLE IF NOT EXISTS music_library "
        query = query + "(year int, artist_name text, album_name text, PRIMARY KEY(year, artist_name))"
        try:
            session.execute(query)
        except Exception as e:
            print(e)

        query1 = "CREATE TABLE IF NOT EXISTS artist_library "
        query1 = query1 + "(artist_name text, album_name text, year int, PRIMARY KEY(artist_name, album_name))"
        try:
            session.execute(query1)
        except Exception as e:
            print(e)

        query2 = "CREATE TABLE IF NOT EXISTS album_library "
        query2 = query2 + "(album_name text, artist_name text, year int, PRIMARY KEY(album_name, artist_name))"
        try:
            session.execute(query2)
        except Exception as e:
            print(e)
```

1.0.9 TO-DO: Insert data into the tables

```
In [8]: query = "INSERT INTO music_library (year, artist_name, album_name)"
        query = query + " VALUES (%s, %s, %s)"

        query1 = "INSERT INTO artist_library (artist_name, year, album_name)"
```

```

query1 = query1 + " VALUES (%s, %s, %s)"

query2 = "INSERT INTO album_library (album_name, artist_name, year)"
query2 = query2 + " VALUES (%s, %s, %s)"

try:
    session.execute(query, (1970, "The Beatles", "Let it Be"))
except Exception as e:
    print(e)

try:
    session.execute(query, (1965, "The Beatles", "Rubber Soul"))
except Exception as e:
    print(e)

try:
    session.execute(query, (1965, "The Who", "My Generation"))
except Exception as e:
    print(e)

try:
    session.execute(query, (1966, "The Monkees", "The Monkees"))
except Exception as e:
    print(e)

try:
    session.execute(query, (1970, "The Carpenters", "Close To You"))
except Exception as e:
    print(e)

try:
    session.execute(query1, ("The Beatles", 1970, "Let it Be"))
except Exception as e:
    print(e)

try:
    session.execute(query1, ("The Beatles", 1965, "Rubber Soul"))
except Exception as e:
    print(e)

try:
    session.execute(query1, ("The Who", 1965, "My Generation"))
except Exception as e:
    print(e)

try:
    session.execute(query1, ("The Monkees", 1966, "The Monkees"))
except Exception as e:

```

```

        print(e)

    try:
        session.execute(query1, ("The Carpenters", 1970, "Close To You"))
    except Exception as e:
        print(e)

    try:
        session.execute(query2, ("Let it Be", "The Beatles", 1970))
    except Exception as e:
        print(e)

    try:
        session.execute(query2, ("Rubber Soul", "The Beatles", 1965))
    except Exception as e:
        print(e)

    try:
        session.execute(query2, ("My Generation", "The Who", 1965))
    except Exception as e:
        print(e)

    try:
        session.execute(query2, ("The Monkees", "The Monkees", 1966))
    except Exception as e:
        print(e)

    try:
        session.execute(query2, ("Close To You", "The Carpenters", 1970))
    except Exception as e:
        print(e)

```

This might have felt unnatural to insert duplicate data into the tables. If I just normalized these tables, I wouldn't have to have extra copies! While this is true, remember there are no JOINS in Apache Cassandra. For the benefit of high availability and scalability, denormalization must be how this is done.

1.0.10 TO-DO: Validate the Data Model

```

In [9]: query = "select * from music_library WHERE YEAR=1970"
        try:
            rows = session.execute(query)
        except Exception as e:
            print(e)

        for row in rows:
            print (row.year, row.artist_name, row.album_name)

```

1970 The Beatles Let it Be

1970 The Carpenters Close To You

1.0.11 Your output should be:

1970 The Beatles Let it Be 1970 The Carpenters Close To You

1.0.12 TO-DO: Validate the Data Model

```
In [11]: query = "select * from artist_library WHERE artist_name='The Beatles'"
        try:
            rows = session.execute(query)
        except Exception as e:
            print(e)

        for row in rows:
            print (row.artist_name, row.album_name, row.year)
```

The Beatles Rubber Soul 1965
The Beatles Let it Be 1970

1.0.13 Your output should be:

The Beatles Rubber Soul 1965 The Beatles Let it Be 1970

1.0.14 TO-DO: Validate the Data Model

```
In [12]: query = "select * from album_library WHERE album_name='Close To You'"
        try:
            rows = session.execute(query)
        except Exception as e:
            print(e)

        for row in rows:
            print (row.artist_name, row.year, row.album_name)
```

The Carpenters 1970 Close To You

1.0.15 Your output should be:

The Carpenters 1970 Close To You

1.0.16 And finally close the session and cluster connection

```
In [13]: session.shutdown()
        cluster.shutdown()
```

```
In [ ]:
```